

```
import pandas as pd

df = pd.read_csv('/content/news.csv')
display(df.head())
```

	Unnamed: 0	title	text	label
0	8476	You Can Smell Hillary's Fear	Daniel Greenfield, a Shillman Journalism Fello...	FAKE
1	10294	Watch The Exact Moment Paul Ryan Committed Pol...	Google Pinterest Digg LinkedIn Reddit Stumbleu...	FAKE
2	3608	Kerry to go to Paris in gesture of sympathy	U.S. Secretary of State John F. Kerry said Mon...	REAL
3	10142	Bernie supporters on Twitter erupt in anger ag...	— Kaydee King (@KaydeeKing) November 9, 2016 T...	FAKE
4	875	The Battle of New York: Why This Primary Matters	It's primary day in New York and front-runners...	REAL

```
print(f"Input text column: 'text'")
print(f"Target label column: 'label'")

# Print dataset size (number of rows and columns)
print(f"Dataset size: {df.shape[0]} rows, {df.shape[1]} columns")

# Print class distribution of the target label
print("\nClass distribution of 'label' column:")
display(df['label'].value_counts())
```

```
Input text column: 'text'
Target label column: 'label'
Dataset size: 6335 rows, 4 columns

Class distribution of 'label' column:
  count
  label
  REAL    3171
  FAKE    3164

dtype: int64
```

```
# Download NLTK stopwords if not already downloaded
try:
    stopwords.words('english')
except LookupError:
    nltk.download('stopwords')

# Initialize Porter Stemmer
ps = PorterStemmer()

# Define a preprocessing function
def preprocess_text(text):
    text = text.lower() # Convert to lowercase
    text = re.sub(r'^[a-z\s]', '', text) # Remove punctuation and numbers, keep only letters and spaces
    words = text.split() # Split into words
    words = [word for word in words if word not in stopwords.words('english')]] # Remove stopwords
    # Optional: Stemming
    # words = [ps.stem(word) for word in words]
    return ' '.join(words)

# Apply preprocessing to the 'text' column
df['text'] = df['text'].apply(preprocess_text)

# Display the first 5 samples of the preprocessed text
print("First 5 samples of preprocessed text:")
display(df['text'].head())
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]  Unzipping corpora/stopwords.zip.
```

```
First 5 samples of preprocessed text:
```

	text
0	daniel greenfield shillman journalism fellow f...
1	google pinterest digg linkedin reddit stumbleu...
2	us secretary state john f kerry said monday st...
3	kaydee king kaydeeking november lesson tonight...
4	primary day new york frontrunners hillary clin...

```
dtype: object
```

```
# Initialize TF-IDF Vectorizer
tfidf_vectorizer = TfidfVectorizer(max_features=5000) # Limiting to 5000 features for demonstration
```

```
# Fit and transform the preprocessed text data
X_tfidf = tfidf_vectorizer.fit_transform(df['text'])
```

```
# Print feature matrix shape
print(f"Feature matrix shape: {X_tfidf.shape}")
```

```
# Display sample feature names
print("\nSample feature names (first 20):")
display(tfidf_vectorizer.get_feature_names_out()[:20])
```

```
Feature matrix shape: (6335, 5000)
```

```
Sample feature names (first 20):
```

```
array(['abandon', 'abandoned', 'abc', 'abdullah', 'abedin', 'ability',
       'able', 'abortion', 'abortions', 'abroad', 'absence', 'absolute',
       'absolutely', 'absurd', 'abuse', 'abuses', 'academic', 'academy',
       'accept', 'acceptable'], dtype=object)
```

```
# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(
    X_tfidf, df['label'], test_size=0.2, random_state=42, stratify=df['label']
)
```

```
# Print the shapes of the resulting datasets
print(f"Shape of X_train: {X_train.shape}")
print(f"Shape of X_test: {X_test.shape}")
print(f"Shape of y_train: {y_train.shape}")
print(f"Shape of y_test: {y_test.shape}")
```

```
Shape of X_train: (5068, 5000)
Shape of X_test: (1267, 5000)
Shape of y_train: (5068,)
Shape of y_test: (1267,)
```

## The Assumption of Conditional Independence in Naive Bayes

The Naive Bayes algorithm is a probabilistic classifier based on Bayes' theorem, with a crucial 'naive' assumption: **conditional independence of features**.

Here's what that means:

1. **Bayes' Theorem:** At its core, Naive Bayes uses Bayes' theorem to calculate the probability of a class (e.g., 'FAKE' or 'REAL') given a set of observed features (e.g., the words in a news article):

$$P(\text{Class}|\text{Features}) = \frac{P(\text{Features}|\text{Class}) \cdot P(\text{Class})}{P(\text{Features})}$$

- $P(\text{Class}|\text{Features})$  is the posterior probability: the probability of the class given the features.
- $P(\text{Features}|\text{Class})$  is the likelihood: the probability of observing the features given the class.
- $P(\text{Class})$  is the prior probability: the probability of the class before seeing any features.
- $P(\text{Features})$  is the evidence: the probability of observing the features.

2. **The 'Naive' Assumption:** To make the calculation of  $P(\text{Features}|\text{Class})$  tractable, especially when dealing with many features (like thousands of words in a text classification problem), Naive Bayes assumes that **each feature is conditionally independent of every other feature, given the class**.

In simpler terms, this means that the presence or absence of a particular word in a document (e.g., 'politics') does not affect the presence or absence of another word (e.g., 'economy'), *given that we already know the document's class* (e.g., it's a 'REAL' news article).

Mathematically, for features  $F_1, F_2, \dots, F_n$  and a class  $C$ :

$$P(F_1, F_2, \dots, F_n | C) = P(F_1 | C) \cdot P(F_2 | C) \cdot \dots \cdot P(F_n | C)$$

**3. Why 'Naive'?**: This assumption is called 'naive' because it's often not true in real-world scenarios. For instance, in a news article, the presence of the word 'president' is highly likely to be correlated with the presence of 'government'. These words are not truly independent.

**4. Despite the Assumption**: Despite this unrealistic assumption, Naive Bayes classifiers often perform surprisingly well in practice, especially in text classification. This is because:

- **Simplicity and Efficiency**: It's computationally inexpensive and fast to train.
- **Robustness to Irrelevant Features**: Irrelevant features (words) don't significantly harm the prediction, as they are averaged out.
- **Good for High-Dimensional Data**: It handles high-dimensional data (many features) well, which is common in text data where each unique word can be a feature.

For text classification tasks, **Multinomial Naive Bayes** is often chosen because it is well-suited for discrete counts (like word counts or TF-IDF values, which are essentially normalized counts) and performs effectively in such scenarios.

```
# Initialize and train the Multinomial Naive Bayes model
mnb_model = MultinomialNB()
mnb_model.fit(X_train, y_train)

# Display model parameters (if any are directly accessible or useful)
# For MultinomialNB, parameters are typically learned probabilities, not 'hyperparameters' like in other models.
# We can look at the classes and feature log probabilities.
print("Model classes:", mnb_model.classes_)
print("Number of features:", mnb_model.n_features_in_)
print("Number of training samples per class:", mnb_model.class_count_)
print("Log prior probability for each class:", mnb_model.class_log_prior_)
print("Log probability of features given a class (first 5 features for each class):")
display(pd.DataFrame(mnb_model.feature_log_prob_[:, :5], columns=tfidf_vectorizer.get_feature_names_out()[:5], index=mnb_model.classes_))

Model classes: ['FAKE' 'REAL']
Number of features: 5000
Number of training samples per class: [2531. 2537.]
Log prior probability for each class: [-0.69433178 -0.69196398]
Log probability of features given a class (first 5 features for each class):
      abandon abandoned      abc abdullah      abedin
FAKE -9.477648 -9.313346 -9.010888 -9.745106 -7.959056
REAL -9.290732 -9.320792 -9.015442 -9.311361 -9.474541
```

```
# Import necessary metrics (if not already imported)
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import pandas as pd # For displaying confusion matrix nicely

# Make predictions on the test set
y_pred = mnb_model.predict(X_test)

print("--- Model Evaluation ---")

# 1. Accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")

# 2. Classification Report (includes Precision, Recall, F1-score)
print("\nClassification Report:")
print(classification_report(y_test, y_pred, target_names=['FAKE', 'REAL']))

# 3. Confusion Matrix
print("\nConfusion Matrix:")
cm = confusion_matrix(y_test, y_pred, labels=['FAKE', 'REAL'])
cm_df = pd.DataFrame(cm, index=['Actual FAKE', 'Actual REAL'], columns=['Predicted FAKE', 'Predicted REAL'])
display(cm_df)

# Interpretation of Confusion Matrix
```

```

print("\nInterpretation of Confusion Matrix:")
print(f"True Negatives (Top-Left): {cm[0,0]} - FAKE news correctly predicted as FAKE")
print(f"False Positives (Top-Right): {cm[0,1]} - FAKE news incorrectly predicted as REAL")
print(f"False Negatives (Bottom-Left): {cm[1,0]} - REAL news incorrectly predicted as FAKE")
print(f"True Positives (Bottom-Right): {cm[1,1]} - REAL news correctly predicted as REAL")

--- Model Evaluation ---
Accuracy: 0.8635

```

Classification Report:

	precision	recall	f1-score	support
FAKE	0.85	0.88	0.87	633
REAL	0.88	0.84	0.86	634
accuracy			0.86	1267
macro avg	0.86	0.86	0.86	1267
weighted avg	0.86	0.86	0.86	1267

Confusion Matrix:

	Predicted FAKE	Predicted REAL
Actual FAKE	559	74
Actual REAL	99	535

Interpretation of Confusion Matrix:

True Negatives (Top-Left): 559 - FAKE news correctly predicted as FAKE  
 False Positives (Top-Right): 74 - FAKE news incorrectly predicted as REAL  
 False Negatives (Bottom-Left): 99 - REAL news incorrectly predicted as FAKE  
 True Positives (Bottom-Right): 535 - REAL news correctly predicted as REAL

```

import matplotlib.pyplot as plt
import seaborn as sns

# Visualize the confusion matrix as a heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(cm_df, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.title('Confusion Matrix for Fake News Detection')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()

```



